

Intro to probabilistic data structures

Download this presentation here:



About Me

- Software Engineer
- Currently at Spotify
- Interest in algorithms & DS

Agenda

Probabilistic Data Structures for this classes of problems:

- membership
- cardinality
- frequency

PDS will return approximate values

Membership



Membership use cases

In general when wrong answer do not involve correctness but more work

- Taken username
- Cassandra avoids checking SSTable data files when merging data on disk
- Ad placement: has the user already seen this ad?
- Fraud detection (has the user paid from this location before?)

Membership

Checking if an item belongs to a set

- Bloom Filter
- Counting Bloom Filter
- Cuckoo Filter

TYPE I ERROR: FALSE POSITIVE

TYPE II ERROR: FALSE NEGATIVE

TYPE III ERROR: TRUE POSITIVE FOR
INCORRECT REASONS

TYPE IV ERROR: TRUE NEGATIVE FOR
INCORRECT REASONS

TYPE V ERROR: INCORRECT RESULT WHICH
LEADS YOU TO A CORRECT
CONCLUSION DUE TO
UNRELATED ERRORS

TYPE VI ERROR: CORRECT RESULT WHICH
YOU INTERPRET WRONG

TYPE VII ERROR: INCORRECT RESULT WHICH
PRODUCES A COOL GRAPH

TYPE VIII ERROR: INCORRECT RESULT WHICH
SPARKS FURTHER RESEARCH
AND THE DEVELOPMENT OF
NEW TOOLS WHICH REVEAL
THE FLAW IN THE ORIGINAL
RESULT WHILE PRODUCING
NOVEL CORRECT RESULTS

TYPE IX ERROR: THE RISE OF SKYWALKER

(source: <https://xkcd.com/2303>)

Bloom Filter

Invented by Burton H. Bloom in 1970

Consists of :


- a bit array of size m
- k different hash functions

Membership check :


- if all bits are set

Bloom Filter with $m=16, k=3$


Insert x



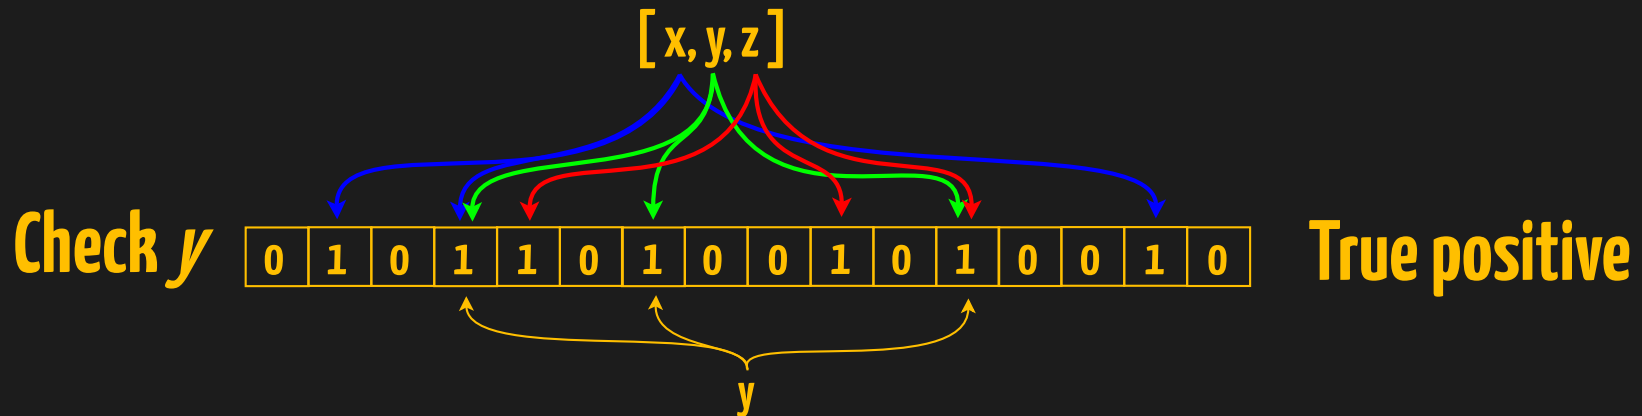
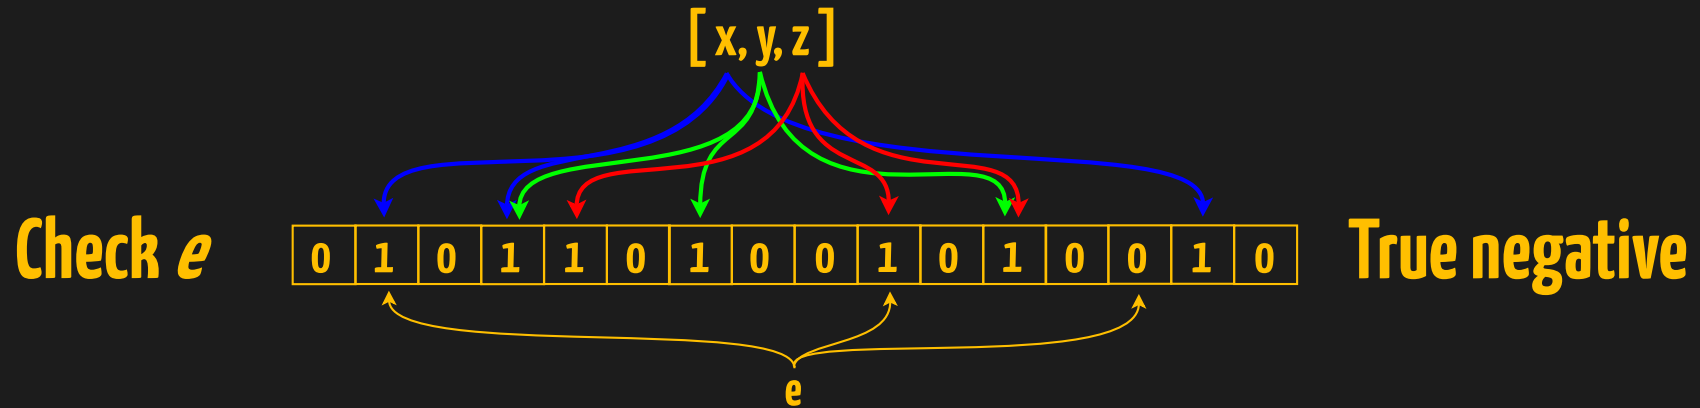
Insert y



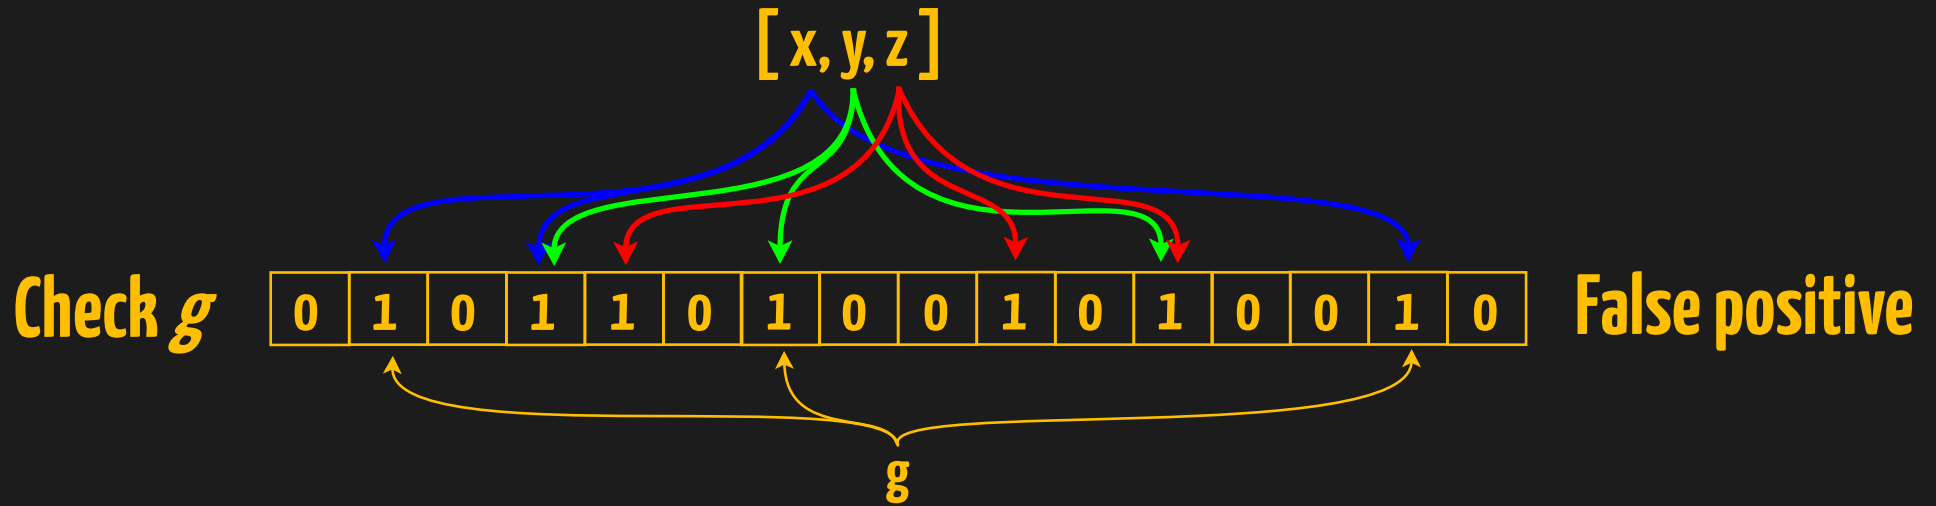
Insert z



Bloom Filter with $m=16, k=3$



Bloom Filter with $m=16, k=3$



Bloom Filter

Approximate optimal size

Given:

- n : the number of expected items
- ε : the wanted false positive rate, where $\varepsilon \in [0, 1]$

Then:

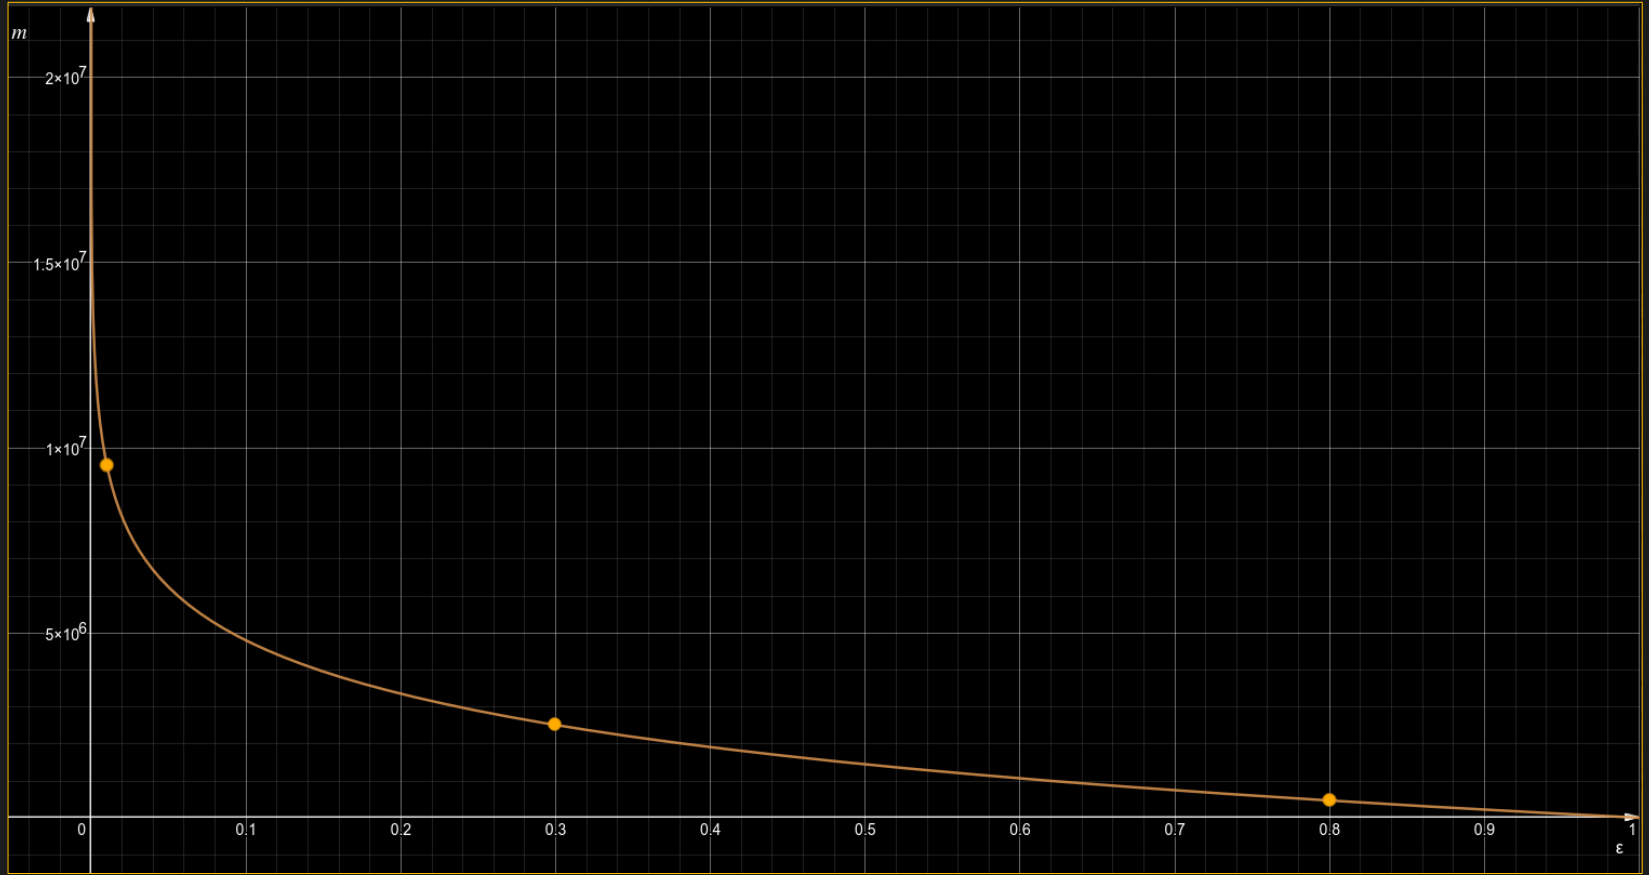
- $m = -\frac{n \cdot \ln(\varepsilon)}{\ln(2)^2}$
- $k = -\log_2(\varepsilon)$

To check 1M items, with 1% error rate:

- $m = -\frac{1,000,000 \cdot -4.6}{0.48} \simeq 9.5 \cdot 10^6 \simeq 1.2 \text{ MB}$
 - $k \simeq 6.6 = 7$
- We have a capped size!

Bloom Filter

Given $n = 1\text{M}$



$$\epsilon = 0.01 \Rightarrow m = 9.5\text{M}$$

$$\epsilon = 0.3 \Rightarrow m = 2.5\text{M}$$

$$\epsilon = 0.8 \Rightarrow m = 460\text{K}$$

Bloom Filter Code

```
01 class StringBloomFilter(expectedSize: Int, errorRate: Double) {
02
03     private val m: Int = -(expectedSize * ln(errorRate) / ln2squared).toInt() // bitset size
04     private val k = ceil(-log2(errorRate)).toInt() // number of hash functions
05     private val bitSet = BitSet(m)
06     private val hashers = IntStream
07         .rangeClosed(1, k)
08         .mapToObj { n → Hasher(primes[n + 4], primes[3 * n + 3]) }
09         .toList()
10
11     fun contains(item: String)=hashers.all { hasher→bitSet.get(abs(hasher.hashCode(item)) % m) }
12
13     fun add(item: String)=hashers.forEach { hasher→bitSet.set(abs(hasher.hashCode(item)) % m, true)}
14
15     private class Hasher(private val base: Int, private val multiplier: Int) {
16         fun hashCode(value: String): Int = value
17             .map { it.code }
18             .fold(base) { acc, curr → acc * multiplier + curr }
19     }
20 }
```

Bloom Filter



Fun Fact

Fruit flies olfactory neural circuit evolved a variant of a Bloom filter to assess the novelty of odors!


It adds two additional features:

- **based on similarity to previously experienced odors**
- **time elapsed since the odor was last experienced**




(source: <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC6304992/>)

Bloom Filter

If we remove an item from the set, we have false negatives!



Bloom Filter with $M=16, k=3$



False negative!

Counting Bloom Filter

Consists of:

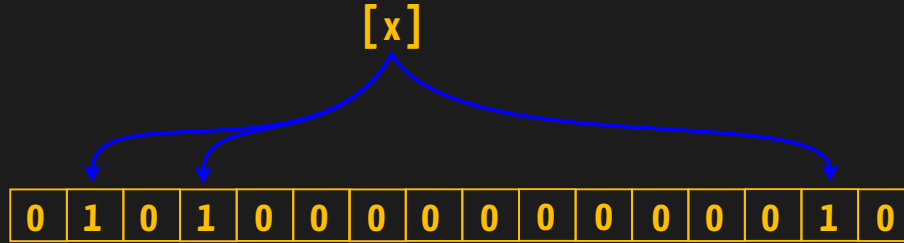
- a count array of size m
- k different hash functions

Membership check:

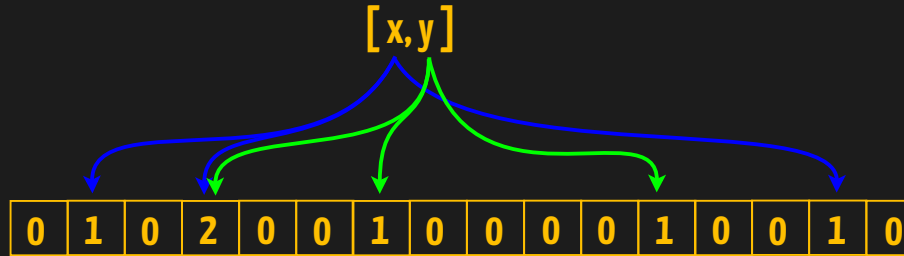
- if all counts are > 0

Counting Bloom Filter

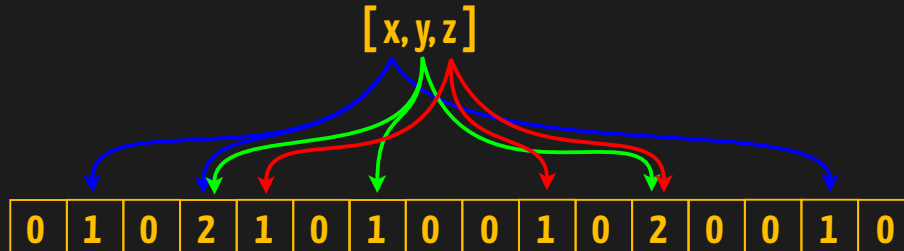
Insert x



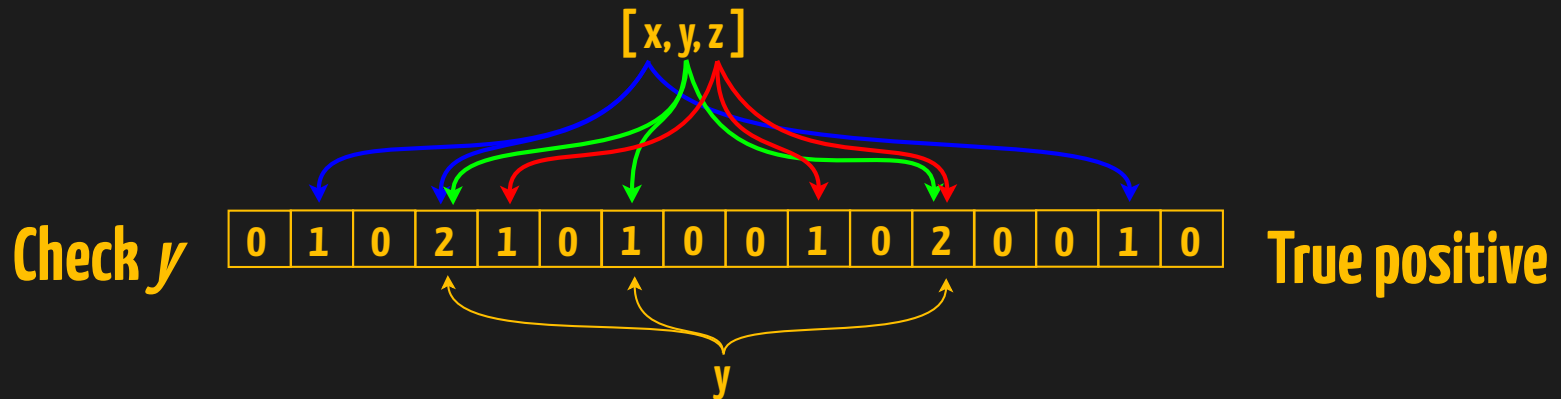
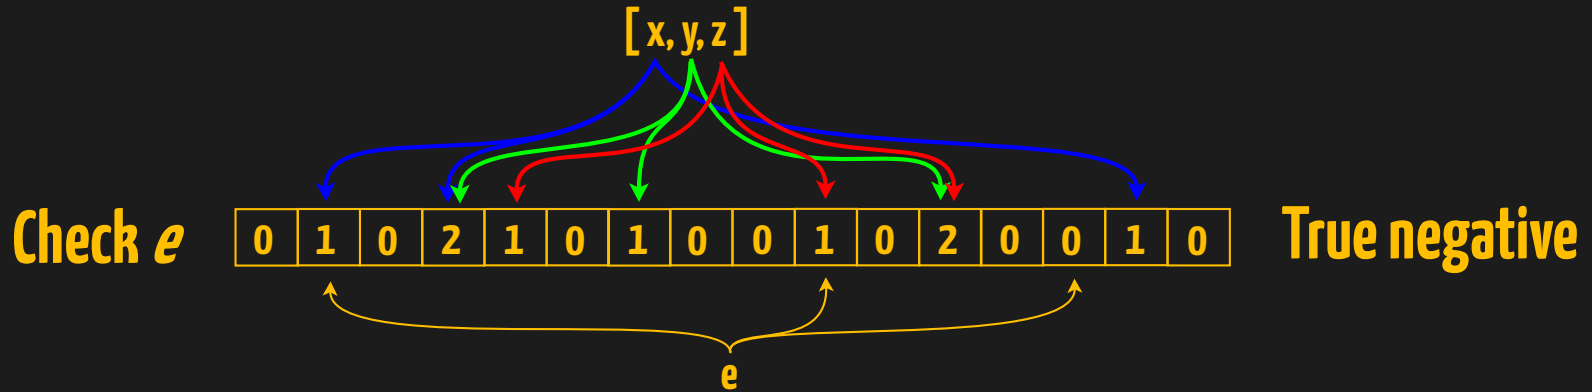
Insert y



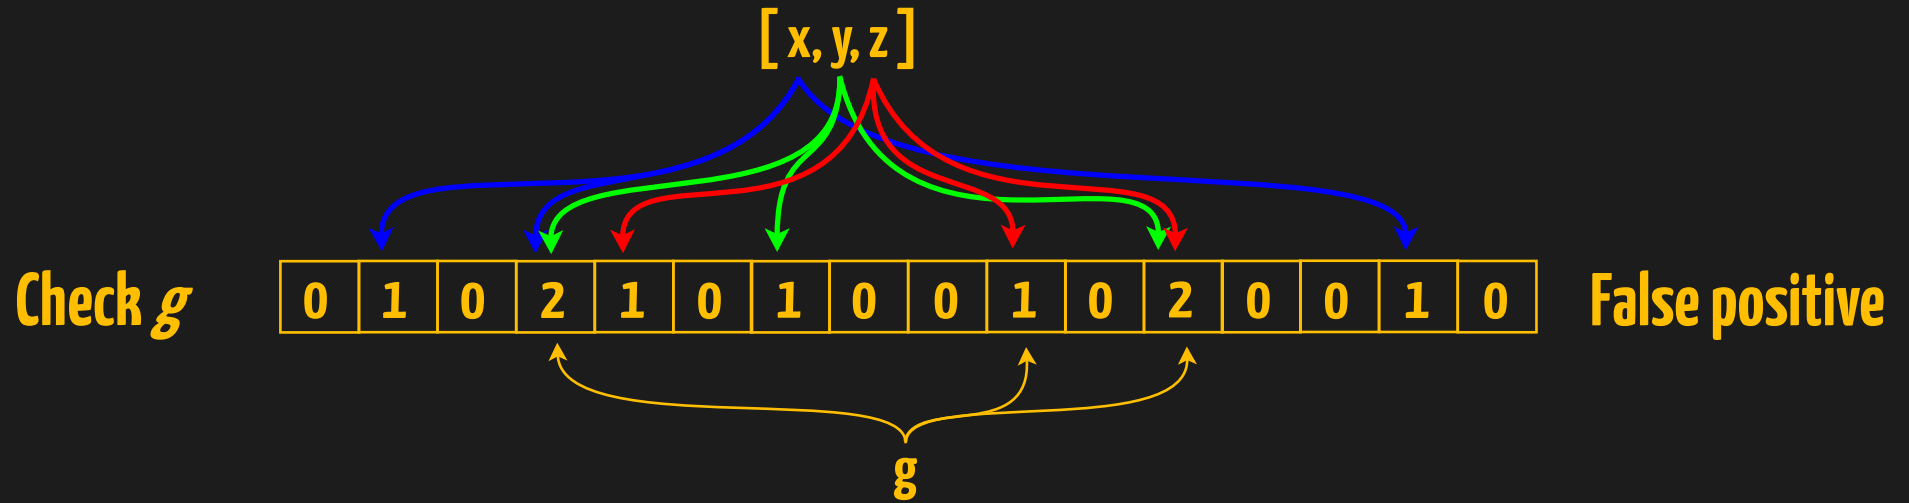
Insert z



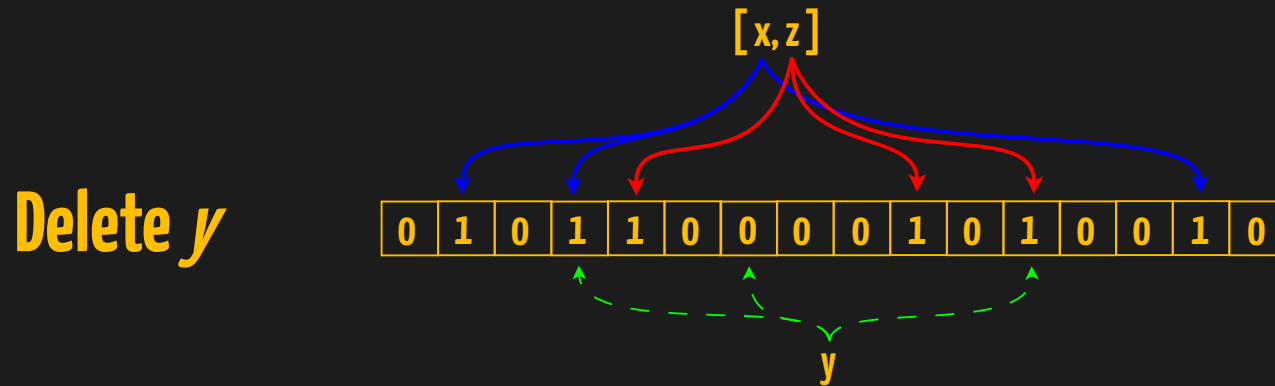
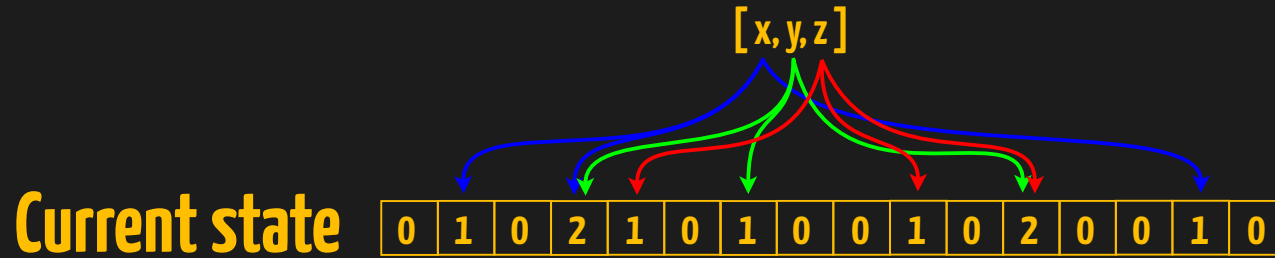
Counting Bloom Filter



Counting Bloom Filter



Counting Bloom Filter



```

01 class StringCountingBloomFilter(expectedSize: Int, errorRate: Double) {
02     private val m: Int = -(expectedSize * ln(errorRate) / ln2squared).toInt()
03     private val k = ceil(-log2(errorRate)).toInt()
04     private val counters = ByteArray(m) { Byte.MIN_VALUE }
05     private val hashers = IntStream
06         .rangeClosed(1, k)
07         .mapToObj { n → Hasher(primes[n + 4], primes[3 * n + 3]) }
08         .toList()
09
10     fun add(item: String)=hashers.forEach { counters[abs(it.hashCode(item)) % m]++ }
11     fun delete(item: String)=hashers.forEach { counters[abs(it.hashCode(item)) % m]-- }
12     fun contains(item: String)=hashers.all { counters[abs(it.hashCode(item)) % m] > MIN_VALUE }
13
14     private class Hasher(private val base: Int, private val mult: Int) {
15         fun hashCode(value: String) = value.map { it.code }
16             .fold(base) { acc, curr → acc * mult + curr }
17     }
18 }
19 }

```


Counting Bloom Filter Code

Counting Bloom Filter

It uses a lot more memory than Bloom Filter:

- 8x with byte
- 16x with short
- 32x with int

It also might overflow



Cuckoo Filter


Described by Fan, Andersen, Kaminsky, and Mitzenmacher in 2014

Cuckoo Hashing:


- 1) Each item to insert has two possible locations given by two hashing functions
- 2^a) If one or both of the two locations are empty, just insert
- 2^b) Else randomly select one of the two locations, insert and kick out the old item
- 3) Insert the old item into the other location: if not empty, repeat this process

Cuckoo Hashing

Choose location for x




Insert and keep track




Cuckoo Hashing

Choose location for y




Insert and keep track




Cuckoo Hashing

Choose location for z

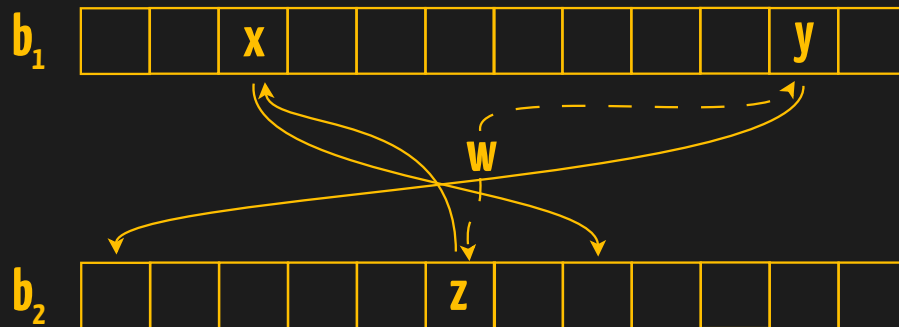


Insert and keep track




Cuckoo Hashing

Choose location for w

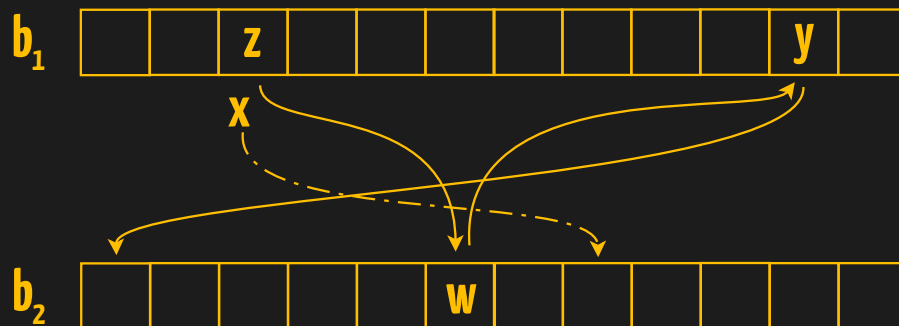


Randomly choose bucket #2
and kick out z !

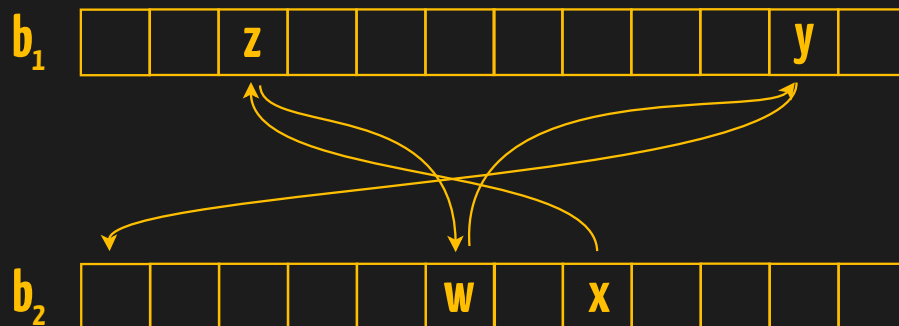


Cuckoo Hashing

z kicks out x



Insert x into its
alternative location




Cuckoo Hashing

- $O(1)$ lookup runtime complexity: each item will be in one of the two locations
- Insertion can fail if there's a loop in items to be kicked out
- Incredibly enough, amortized runtime complexity of insertion is $O(1)$

Cuckoo Filter

Based on Cuckoo Hashing

- to save memory instead of storing the whole item we store a fingerprint (f bit)
- how can we get the alternate location if we don't have the kicked out item?



- create a hash that keeps track of fingerprint

Cuckoo Filter

- a fingerprint function returning a f -bit value

- two hash functions
$$\begin{cases} h_1(x) = \text{hash}(x) \\ h_2(x) = h_1(x) \oplus \text{hash}(\text{fingerprint}(x)) \end{cases}$$
$$\Downarrow$$
$$h_1(x) = h_2(x) \oplus \text{hash}(\text{fingerprint}(x))$$

Cuckoo Filter

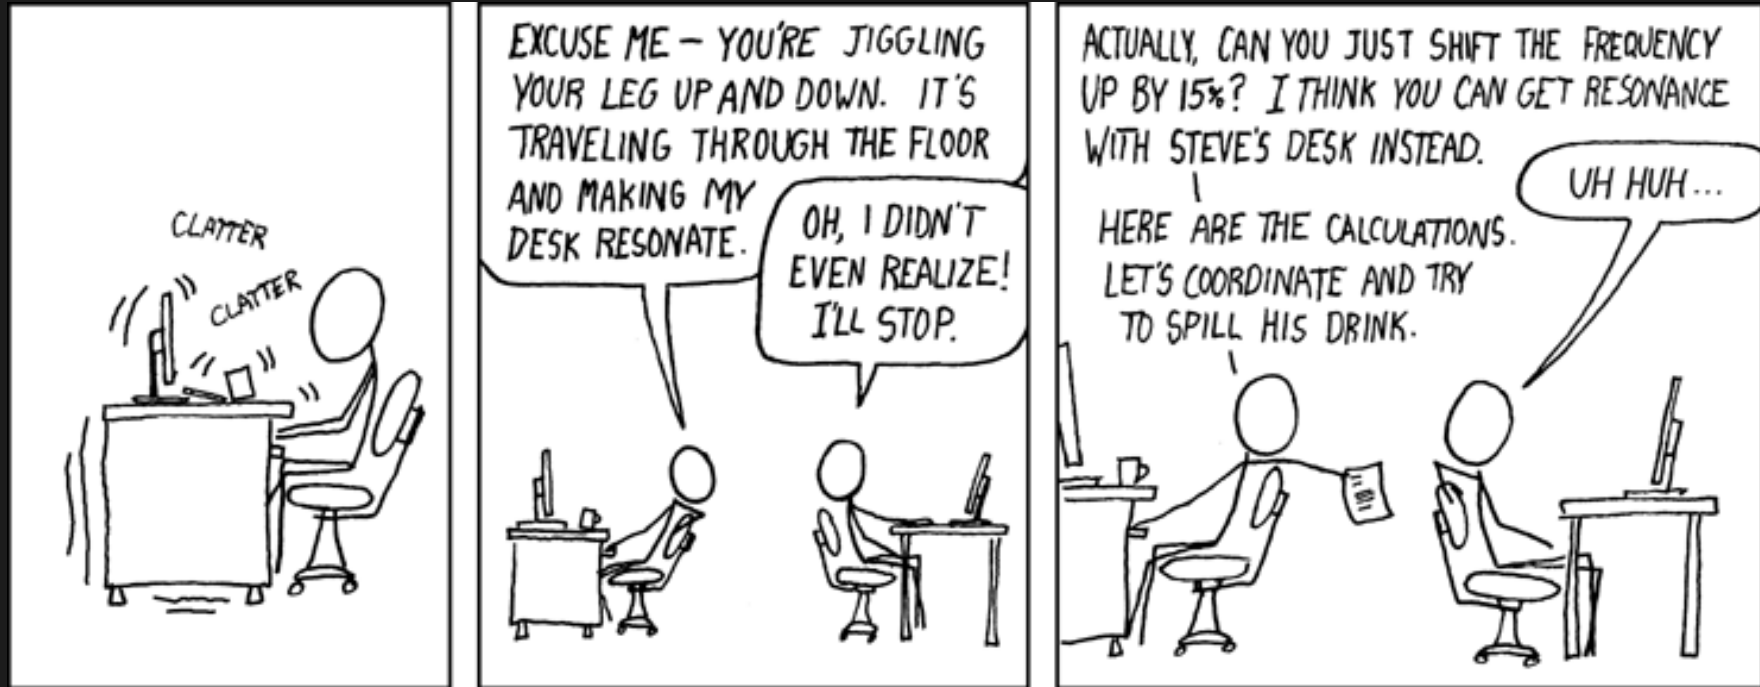
Cuckoo filter requires $\frac{\log_2\left(\frac{1}{\epsilon}\right) + 1 + \log_2(b)}{\alpha}$ bits per key

where α is the load factor of the cuckoo hash table

Cuckoo Filter Code

```
01 fun insert(item: String): Boolean {
02     var f = fingerprint(item)
03     val i1 = hash(item)
04     if (bucket[i1] == EMPTY) {
05         bucket[i1] = f
06         return true
07     }
08     val i2 = i1 xor hash(f)
09     if (bucket[i2] == EMPTY) {
10         bucket[i2] = f
11         return true
12     }
13
14     // relocate existing items
15     var i = if (Random.nextBoolean()) i1 else i2
16     for (n in 0..
```

Frequency



(source: <https://xkcd.com/228/>)

Frequency

Counting the number of times an item appears in a stream

- Count-Min Sketch

Frequency Use Cases

- Unique users of a service
- Top k items (e.g. top players in online gaming)
- Network traffic analysis

Count-Min Sketch

Invented by G. Cormode et al in 2005

- d hash functions (pairwise independent)
 - array *count* of size $d \times w$ for counting the items
 - $d = \ln\left(\frac{1}{\delta}\right)$
 - $w = \left(\frac{e}{\epsilon}\right)$
- where ϵ is the error rate and δ its probability

If we want an error of at most 0.1% (of the sum of all frequencies) with 99.9% certainty, then:

$$w = \left(\frac{e}{0.001}\right) \simeq 2,718 \quad d = \ln\left(\frac{1}{0.001}\right) \simeq 6.9 = 7$$

Using 32 bit counters, $\text{sizeof}(\textit{count}) = w \cdot d \cdot 4 \simeq 76 \text{ KB}$

Count-Min Sketch

$d = 4, w = 16$

Initial state

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

Hash Function #1

Hash Function #2

Hash Function #3

Hash Function #4

Insert x

$$h_1(x) = 2$$

$$h_2(x) = 13$$

$$h_3(x) = 12$$

$$h_4(x) = 7$$

0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0
0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0

Hash Function #1

Hash Function #2

Hash Function #3

Hash Function #4

Count-Min Sketch

$d = 4, w = 16$

Insert y

$$h_1(y) = 6$$

$$h_2(y) = 13$$

$$h_3(y) = 8$$

$$h_4(y) = 2$$

0	0	1	0	0	0	1	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	2	0	0
0	0	0	0	0	0	0	0	1	0	0	0	1	0	0	0
0	0	1	0	0	0	0	1	0	0	0	0	0	0	0	0

Hash Function #1

Hash Function #2

Hash Function #3

Hash Function #4

Insert x

$$h_1(x) = 2$$

$$h_2(x) = 13$$

$$h_3(x) = 12$$

$$h_4(x) = 7$$

0	0	2	0	0	0	1	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	3	0	0
0	0	0	0	0	0	0	0	1	0	0	0	2	0	0	0
0	0	1	0	0	0	0	2	0	0	0	0	0	0	0	0

Hash Function #1

Hash Function #2

Hash Function #3

Hash Function #4

Count-Min Sketch

$d = 4, w = 16$

0	0	2	0	0	0	1	0	0	0	0	0	0	0	0	0	Hash Function #1
0	0	0	0	0	0	0	0	0	0	0	0	0	3	0	0	Hash Function #2
0	0	0	0	0	0	0	0	1	0	0	0	2	0	0	0	Hash Function #3
0	0	1	0	0	0	0	2	0	0	0	0	0	0	0	0	Hash Function #4

$$\text{Count}(x) = \min(\text{count}[0][h_1(x)], \\ \text{count}[1][h_2(x)], \\ \text{count}[2][h_3(x)], \\ \text{count}[3][h_4(x)])$$

$$= \min(2, 3, 2, 2)$$

$$= 2$$

$$\text{Count}(y) = \min(\text{count}[0][h_1(y)], \\ \text{count}[1][h_2(y)], \\ \text{count}[2][h_3(y)], \\ \text{count}[3][h_4(y)])$$

$$= \min(1, 3, 1, 1)$$

$$= 1$$

Count-Min Sketch

If there are collisions on hashes we get higher counts

Count-Min Sketch Code

```
01 class CountMinSketch(d: Int, private val w: Int) {
02     private val count = Array(d) { LongArray(w) }
03     private val hashers = IntStream
04         .rangeClosed(1, h)
05         .mapToObj { n → Hasher(w, Primes.primes[n * 17 + 4], Primes.primes[3 * n + 3]) }
06         .toList()
07
08     fun add(item: String) = hashers.forEachIndexed { idx, hasher →
09         count[idx][hasher.hashCode(item)]++
10     }
11
12     fun count(item: String) = hashers
13         .mapIndexed { idx, hasher → count[idx][hasher.hashCode(item)] }
14         .min()
15
16     private class Hasher(val size: Int, val base: Long, val multiplier: Long) {
17         fun hashCode(value: String) = value.map { it.code }
18             .fold(base) { acc, curr → acc * multiplier + curr }
19             .mod(size)
20     }
21 }
```

Cardinality



Cardinality

Counting the number of distinct items in a collection

- HyperLogLog

Cardinality Use Cases

- **BigQuery:** APPROX_COUNT_DISTINCT, APPROX_QUANTILES, APPROX_TOP_COUNT, APPROX_TOP_SUM
- **Snowflake:** HLL, HLL_ACCUMULATE, HLL_COMBINE, HLL_ESTIMATE, HLL_ACCUMULATE, HLL_COMBINE
- **Number of users in search engines (where ads are paid per user)**
- **Materialized views in Data Warehouses**

HyperLogLog

Flajolet et al. 2007

Binary representation of random numbers:

- $\frac{1}{2}$ start with “0”
 - $\frac{1}{2}$ start with “1”
 - $\frac{1}{4}$ start with “00”
 - $\frac{1}{4}$ start with “01”
 - $\frac{1}{4}$ start with “10”
 - $\frac{1}{4}$ start with “11”
 - $\frac{1}{8}$ start with “000”
 - $\frac{1}{8}$ start with “001”
 - $\frac{1}{8}$ start with “010”
 - $\frac{1}{8}$ start with “011”
 - $\frac{1}{8}$ start with “111”
 - $\frac{1}{8}$ start with “101”
 - $\frac{1}{8}$ start with “110”
 - $\frac{1}{8}$ start with “111”
- General rule: $p(\text{first } k \text{ bits}) = 2^{-k}$

HyperLogLog

Underlying idea: let's imagine we have a big set of items hashed to numbers; if we add to the set the hash of a new item starting with "0000" we can say that it's likely that the set has size $2^4 = 16$.

Caveats:

- only works on big cardinalities
- hash function must return uniformly distributed binary representations

HyperLogLog

Consists of:

- hash function
- array of int


With only one hash function, we risk that one single item can skew the result. We could use multiple hash functions, but that takes extra computation time

HyperLogLog


Since we don't want multiple hash functions, we use only one, but we pretend they're 2^b



HyperLogLog



HyperLogLog



array

0	0	0	0	0	4	0	0	0	0	0	0	1	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

HyperLogLog



array

0	0	0	0	0	4	0	0	0	0	0	0	1	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

HyperLogLog

Get the count:

- compute a normalized harmonic mean $Z = \left(\sum_{i=1}^m 2^{-M[i]} \right)^{-1}$ of the array values
- return the harmonic mean with the correction factor α applied

$$\alpha = \begin{cases} \alpha_{16} = 0.673 \\ \alpha_{32} = 0.697 \\ \alpha_{64} = 0.709 \\ \alpha_m = \frac{0.7213}{1 + \frac{1.079}{m}} \text{ for } m \geq 128 \end{cases}$$

$$\text{Error rate} = \frac{1.04}{\sqrt{m}}$$

$$m = 16 \text{ (4 bits)} \rightarrow 0.065\%$$

HyperLogLog Code

```
01 class HyperLogLog {
02     private val p: Int = 4
03     private val m: Int = 2.0.pow(p.toDouble()).toInt()
04     private val alpha = 0.673
05     private val buckets = IntArray(m)
06
07     fun add(item: String) {
08         val hash = item.hashCode()
09         val index = hash.ushr(32 - p)
10         val leadingZeroes = Integer.numberOfLeadingZeros(hash shl p) + 1
11         buckets[index] = max(buckets[index], leadingZeroes)
12     }
13
14     fun count(): Long {
15         val harmonicMean = buckets.sumOf { 1.0 / (1 shl it) }
16         val estimate = alpha * (m * m).toDouble() / harmonicMean
17         return if (estimate ≤ 5.0 / 2.0 * m) {
18             (m * ln(m.toDouble() / estimate)).toLong()
19         } else {
20             estimate.toLong()
21         }
22     }
23 }
```

Links

- **Bloom filter:** <https://dl.acm.org/doi/pdf/10.1145/362686.362692>
- **Cuckoo Filter:** <https://www.cs.cmu.edu/~dga/papers/cuckoo-conext2014.pdf>
- **HyperLogLog:** <https://algo.inria.fr/flajolet/Publications/FlFuGaMe07.pdf>
- **Count-min Sketch:** <http://dimacs.rutgers.edu/~graham/pubs/papers/cm-full.pdf>
- **Code for this presentation:** <https://github.com/andreaiacono/TalkProbabilisticDataStructures>

Questions?